

SPRINT 2 · THE GHOSTWRITER

Context Is King

Give an Agent Memory

MS2.1 · Tuesday

"Drafts that sound like you, not a chatbot."

-30 min · Context windows · Tone profiles · Thread history

Sprint 2 — The Ghostwriter

Last week: Read & Triage (The Switchboard)

This week: Draft in Your Voice (The Ghostwriter)

Next week: Take Action (The Operator)

Tuesday (Today)

Context Is King — give the agent memory

Wednesday

The Draft Machine — replies that don't sound robotic

Thursday

Human in the Loop — approve/edit/reject gate

Saturday

Sprint Build: The Draft Desk — 3 real threads drafted

The Problem — Why AI Drafts Suck

- You ask ChatGPT to "reply to this email"
- It writes: "I hope this email finds you well..."
- It sounds generic, corporate, not like you at all

Why? Because it has ZERO context:

- It doesn't know the thread history (who said what)
- It doesn't know your past replies (your voice)
- It doesn't know your role or relationship with the sender
- It doesn't know your preferences (short? formal? casual?)

Today we fix this. We give the agent YOUR context.

What Is a Context Window?

The context window = everything the LLM "sees" before generating.

SYSTEM PROMPT → Your tone profile, role, rules	
CONTEXT → Full thread history (who said what, when) → Your past replies in similar threads → Relationship with sender (boss? client?)	
USER PROMPT → "Draft a reply to this thread"	

More relevant context = better drafts. Garbage in = garbage out.

The 3 Layers of Context We'll Build

1. Thread History

The full email chain — who wrote what, in order. The agent reads the conversation.

2. Tone / Persona Profile

How YOU write — sentence length, formality, sign-off style, common phrases.

3. Past Replies (Few-Shot Examples)

2-3 examples of YOUR actual replies — the model mimics your real writing style.

Layer 1: Thread History

From Sprint 1, you already have this! Your Gmail MCP server gives:

Thread: "Q3 Budget Review"

From: boss@company.com (June 10, 2:30 PM)

Can you review the Q3 budget doc and share your thoughts by EOD tomorrow? Specifically the marketing allocation – I think we're overspending on paid ads.

From: you@company.com (June 10, 3:15 PM)

Got it, will review tonight. Quick question – should I loop in the marketing team lead or keep it between us?

From: boss@company.com (June 11, 9:00 AM)

Keep it between us for now. Let me know what you find.

Layer 2: Tone / Persona Profile

A JSON file that describes HOW you write:

```
{
  "name": "Rahul",
  "role": "Senior Product Manager",
  "tone": "professional but warm",
  "formality": "semi-formal (not stiff, not casual)",
  "sentence_length": "short to medium, gets to the point",
  "greeting_style": "Hi [Name]," or just start talking",
  "sign_off": "Best, Rahul" or "Thanks, R",
  "quirks": [
    "Uses dashes – like this – for asides",
    "Asks one clear question per email",
    "Never uses 'I hope this finds you well'",
    "Keeps replies under 5 sentences when possible"
  ]
}
```

Layer 3: Past Replies (Few-Shot Examples)

- Pick 2–3 emails YOU actually sent (from your Sent folder)
- These become examples for the model — "write like this"
- The model learns your voice from real data, not descriptions

Example from your sent folder:

Subject: Re: Design review feedback

Hi Priya,

Reviewed the mockups – love the direction on the onboarding flow. Two things:

1. The CTA on screen 3 feels buried. Can we make it more prominent?
2. Font size on mobile looks tight. Worth testing on a real device.

Otherwise good to go. Ship it.

Putting It All Together — The Prompt Structure

SYSTEM:

You are drafting email replies for {name}, a {role}.

Tone: {tone}. Formality: {formality}.

Rules: {quirks as bullet points}

Here are examples of how {name} writes:

[Example 1]

[Example 2]

USER:

Here's the email thread:

[Full thread history — newest at bottom]

Draft a reply to the latest message.

Match the tone and style from the examples.

Keep it concise. One clear ask or response.

ASSISTANT:

→ The draft that sounds like YOU

Let's Build It!

Create your tone profile + build the context assembler

Thread history + Tone profile + Past replies = Perfect draft

Today's Output

1. `tone_profile.json` — YOUR writing style captured
2. `past_replies.json` — 2-3 real emails you've sent
3. `context_builder.py` — assembles the full prompt:
 - Loads tone profile
 - Formats thread history
 - Injects past reply examples
 - Returns a ready-to-send prompt for the LLM

Tomorrow: we feed this to Groq and get actual drafts back.

Recap + Assignment

Today we built:

- ✓ Understood the 3 context layers
- ✓ Created your personal tone profile
- ✓ Collected past reply examples
- ✓ Built a context assembler function

Your assignment:

- Write YOUR real tone_profile.json
- Pull 2-3 emails from YOUR sent folder
- Run the context builder on a real thread
- Share the assembled prompt in Slack

Tomorrow: "The Draft Machine" — we connect Groq and get real replies.

See You Tomorrow!

Wednesday: "The Draft Machine" — AI replies that don't sound robotic

#MyAIChiefOfStaff